CS 5/7320 Artificial Intelligence

Search with Uncertainty

AIMA Chapters 4.3-4.5

Slides by Michael Hahsler



Types of uncertainty we consider for now*



Nondeterministic Actions:

Outcome of an action in a state is uncertain.



No observations:

Sensorless problem



Partially observable environments: The agent does not know in what state it is



Unknown environments: Online search

* we will quantify uncertainty with probabilities later.

Consequence of Uncertainty

 Remember: The solution for the known maze was a fixed sequence of actions from start to goal.

 With uncertainty: Solution is not a precomputed sequence, but a

conditional plan (a strategy or policy)

that depends on percepts.



Nondeterministic Actions

Outcome of actions in the environment is nondeterministic = transition model need to describe uncertainty

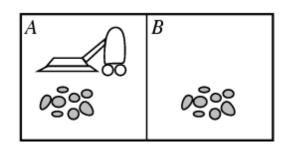
Example transition:

 $Results(s_1, a) = \{s_2, s_4, s_5\}$

i.e., action a in s_1 can lead to one of several states.

The set of possible states $b = \{s_2, s_4, s_5\}$ is called a **belief state** of the agent.

Example: Erratic Vacuum World

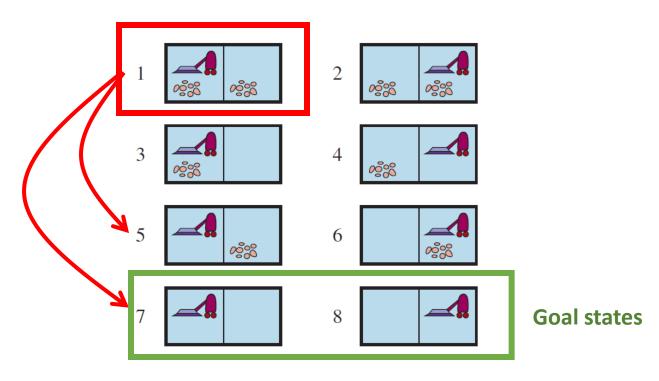


Regular vacuum world, but the action 'suck' is more powerful and nondeterministic:

- a) On a dirty square: cleans the square and sometimes cleans dirt on adjacent squares as well.
- **b)** On a clean square: sometimes deposits some dirt on the square.

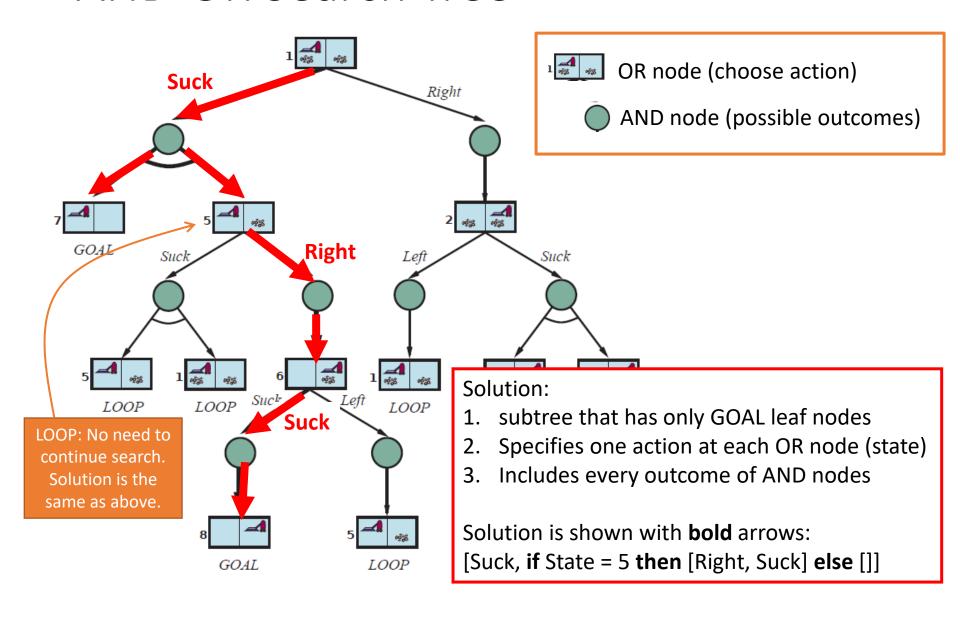
Example: Erratic Vacuum World

 $Results(1, Suck) = \{5, 7\}$



We need a conditional plan
[Suck, if State = 5 then [Right, Suck] else []]

AND-OR Search Tree



AND-OR Recursive DFS Algorithm

= nested If-then-else statements

```
function AND-OR-SEARCH(problem) returns a conditional plan, or failure
  return OR-SEARCH(problem, problem.INITIAL, [])
                                                              path is only maintained for cycle checking!
function OR-SEARCH(problem, state, path) returns a conditional plan, or failure
  if problem.IS-GOAL(state) then return the empty plan
  if IS-CYCLE(path) then return failure
                                                         // don't follow loops using path.
  for each action in problem.ACTIONS(state) do
                                                         // try all possible actions
      plan \leftarrow AND\text{-SEARCH}(problem, RESULTS(state, action), [state] + path])
      if plan \neq failure then return [action] + plan
  return failure
function AND-SEARCH(problem, states, path) returns a conditional plan, or failure
                                                         // try all possible outcomes, none can fail!
  for each s_i in states do
                                                          // (= belief state)
      plan_i \leftarrow \text{OR-SEARCH}(problem, s_i, path)
      if plan_i = failure then return failure
  return [if s_1 then plan_1 else if s_2 then plan_2 else ... if s_{n-1} then plan_{n-1} else plan_n]
```

BFS and A* search can also be used to search an AND-OR tree.

Search with no Observations

No Observations

Sensorless problem = conformant problem

 Example: Doctor prescribes a broad-band antibiotic instead of performing time-consuming blood work for a more specific antibiotic. This saves time and money.

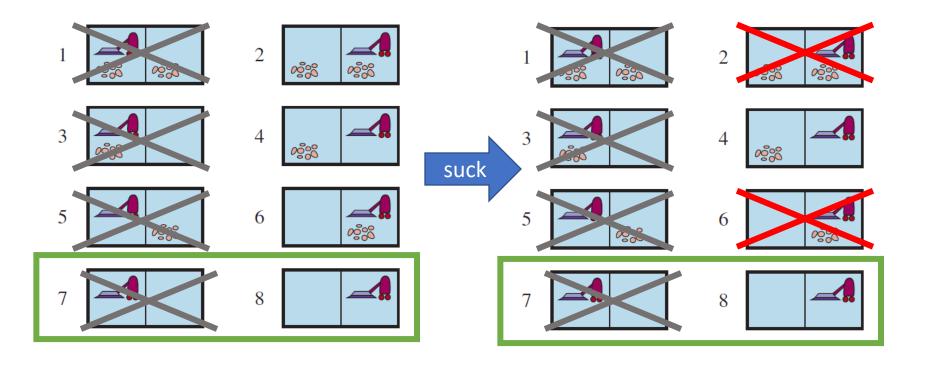
• Basic idea: Find a solution (a sequence of actions) that works from any state.

Actions to Coerce the World into States

- Actions can reduce the number of possible states.
- **Example**: Deterministic vacuum world. Agent does not know its position and the dirt distribution.

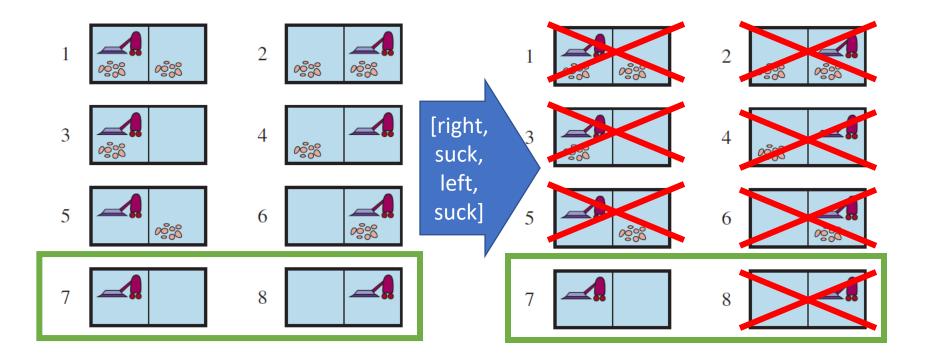
Actions to Coerce the World into States

- Actions can reduce the number of possible states.
- **Example**: Deterministic vacuum world. Agent does not know its position and the dirt distribution.



Actions to Coerce the World into States

- The action sequence [right, suck, left, suck] coerces the world into the goal state 7.
- This works from any initial state and there is no need for a conditional plan!



Find the Solution Sequence

Use regular search (DFS, BFS, A*) for the problem:/

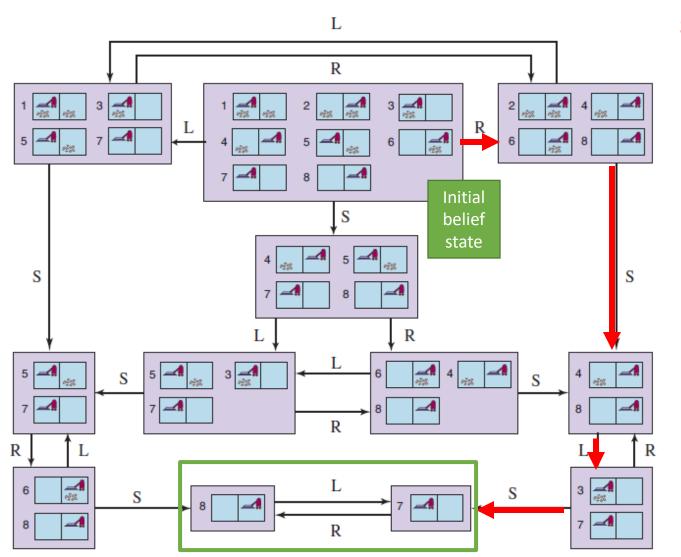
Note: State space size makes this impractical for larger problems!

- States: All belief states (=powerset \mathcal{P}_s of states of size 2^N for N states)
- Initial state: Often the belief state consisting of all states.
- Actions: Actions of a belief state are the union of the possible actions for all the states it contains.
- Transition model: $b' = Results(b, a) = \{s' : s' = Result(s, a) \text{ and } s \in b\}$
- Goal test: Are all states in the belief state goal states?
- **Simplifying property:** If a belief state (e.g., $b_1 = \{1,2,3,4,5\}$) is solvable (i.e., there is a sequence of actions that coerce all states to only goal states), then belief states that are subsets (e.g., $b_2 = \{2,5\}$) are also solved using the same action sequence. Used to prune the search tree.

Other approach:

 Incremental belief-state search. Generate a solution that works for one state and check if it also works for all other states.

Reachable belief-state space for the deterministic, sensorless vacuum world



Solution Sequence: [right, suck, left, suck]

Size of the belief state space:

$$P_s = 2^N = 2^8 = 256$$

Only as small fraction is reachable.

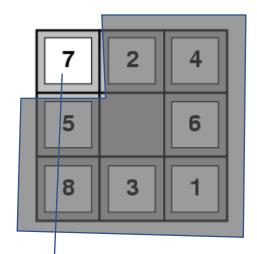


Percepts

- Many problems cannot be solved efficiently without sensing (e.g., 8puzzle).
- We need to be able to at least see one square.

Percept function: Percept(s)

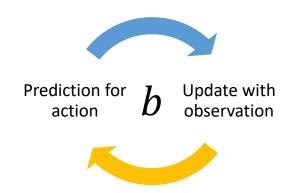
- Fully observable: Percept(s) = s
- Sensorless: Percept(s) = null
- Partially observable: Percept(s) = o

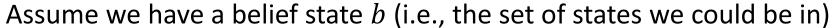


Percept(s) = Tile7

Problem: Many states can produce the same percept!

Prediction and Update





Prediction: Compute new belief state that results from action a.

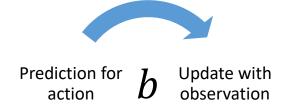
$$\hat{b} = Predict(b, a)$$

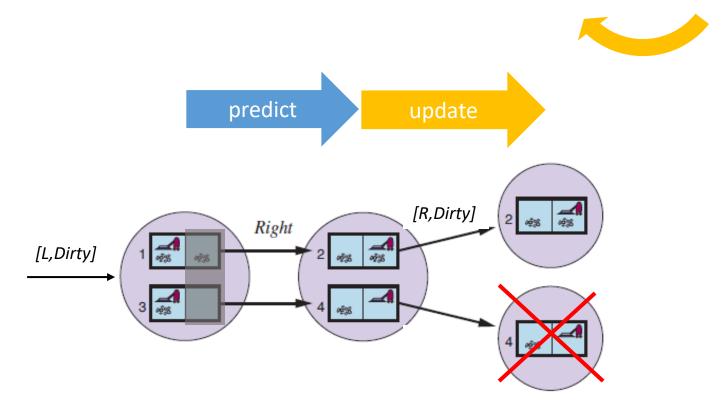
Update: Filter states that are consistent with new observation o.

$$b_o = Update(\hat{b}, o) = \{s : o = Percept(s) \ and \ s \in \hat{b}\}$$

Both steps in one: $b \leftarrow Update(Predict(b, a), o)$

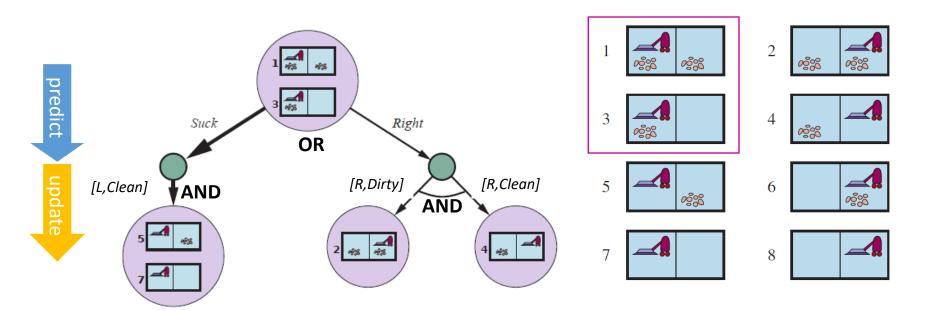
Example: Deterministic local sensing vacuum world





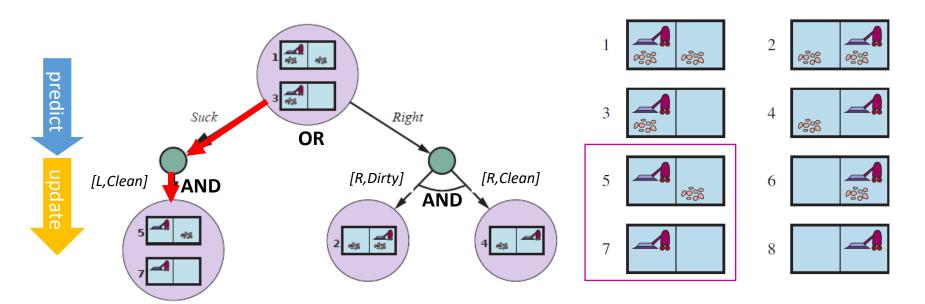
 $b \leftarrow Update(Predict(b, a), o)$ $Update(Predict(\{1,3\}, Right), [R. Dirty]) = \{2\}$

Use an AND-OR tree to create a conditional plan



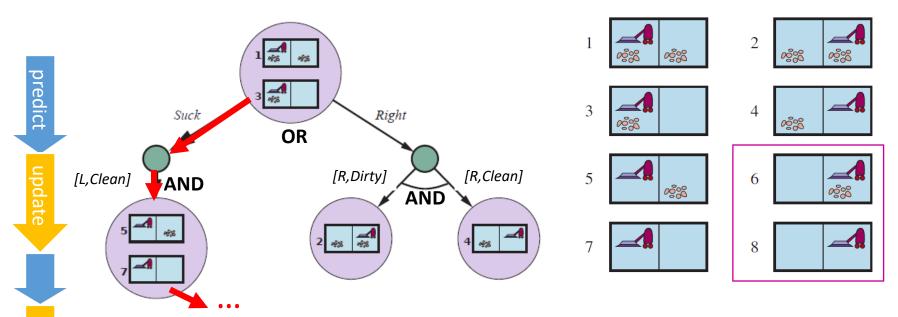
Solution: [Suck, Right, if $b = \{6\}$ then Suck else []]

Use an AND-OR tree to create a conditional plan




Solution: [Suck, Right, if $b = \{6\}$ then Suck else []]

Use an AND-OR tree to create a conditional plan



Solution: [Suck, Right, if $b = \{6\}$ then Suck else []]

Use an AND-OR tree to create a conditional plan



Solution: [Suck, Right, if $b = \{6\}$ then Suck else []]

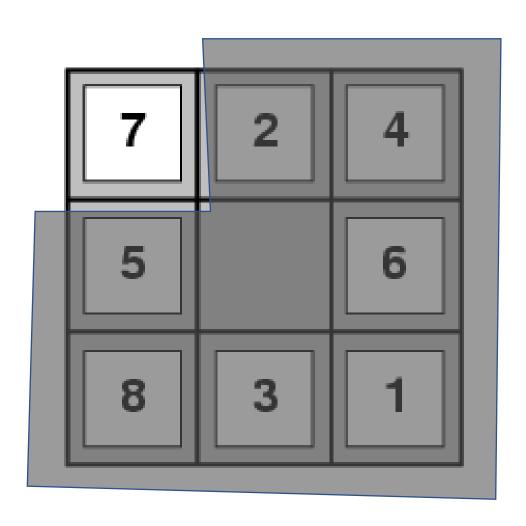
Note: Belief states that are subsets can be used for pruning!

State Estimation and Approximate Belief States

- Agents choose an action and then receive an observation from the environment.
- The agent keep track of its belief state using the following update:

$$b \leftarrow Update(Predict(b, a), o)$$

- This process is often called
 - · monitoring,
 - **filtering**, or
 - state estimation.
- The agent needs to be able to update its belief state following observations in **real time!** For many practical application, there is only time to compute an **approximate belief state!** These approximate methods are outside the scope of this introductory course.

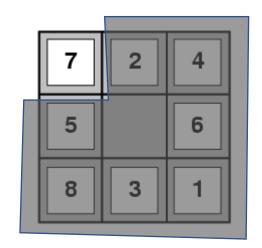


Case Study

Partially Observable 8-Puzzle

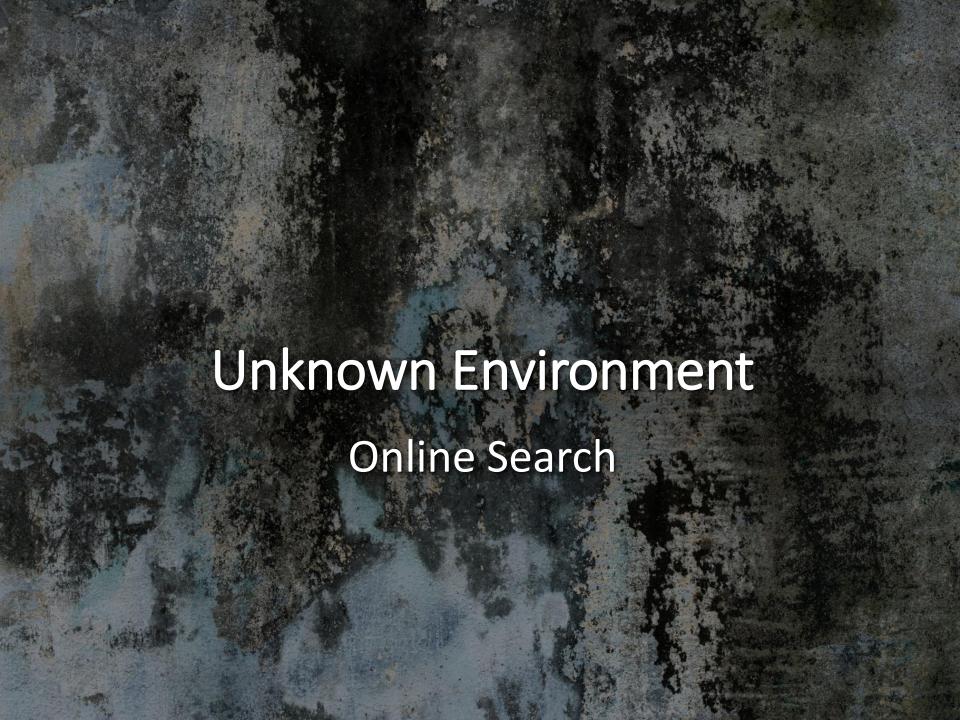
Partially Observable 8-Puzzle

 How do we solve this problem? What are the main steps?



- Give a problem description for each step.
 - States:
 - Initial state:
 - Actions:
 - Transition model:
 - Goal test:
 - Percept function:

What algorithms can be used?



Online Search

- Recall offline search: Create plan using the state space as a model before taking any action. The plan can be a sequence of actions or a conditional plan (for all possible observations).
- Online search explores the real world one action at a time.
 Prediction is replaced by "act" and update by "observe."



- Useful for
 - **Real-time problems**: When offline computation takes too long and there is a penalty for sitting around and thinking.
 - Nondeterministic domain: Deal with what actually happens.
 - **Unknown environment**: The agent needs to explore to map an unknown area (state space). The map is the transition model $f: S \times A \to S$

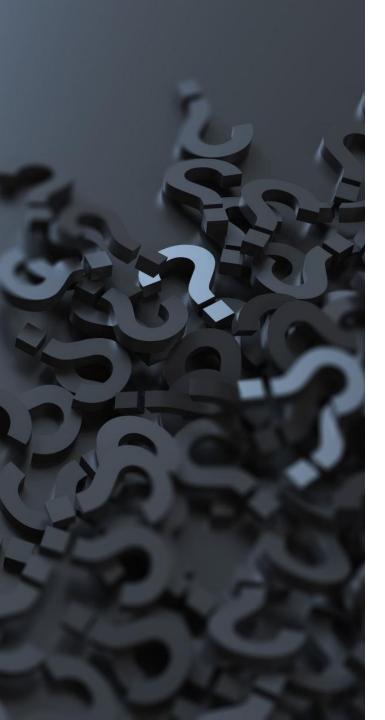
Issues With Online Search

- Knowledge: What does the agent know about the outcome of actions? E.g., does go north and then south lead to the same location?
- Safely explorable state space/world: There are no irreversible actions that cannot be undone (e.g., traps, cliffs). At least the agent does not execute these actions.
- Expanding nodes in local order is more efficient:
 Depth-first search with back tracking

DFS Online Search Algorithm

Environment is deterministic and completely observable (percept(s) = s' after action a), but the transition model (function result) is unknown. This algorithm builds the map result by trying all actions and backtracks when all actions in a state have been explored.

```
function ONLINE-DFS-AGENT(problem, s') returns an action
                                                                             Learns results and
               s, a, the previous state and action, initially null
                                                                                  uses it for
  persistent: result, a table mapping (s, a) to s', initially empty
                                                                                backtracking
               untried, a table mapping s to a list of untried actions
               unbacktracked, a table mapping s to a list of states never backtracked to
  if problem.IS-GOAL(s') then return stop
  if s' is a new state (not in untried) then untried [s'] \leftarrow problem.ACTIONS(s')
  if s is not null then
      result[s, a] \leftarrow s'
      add s to the front of unbacktracked[s'] // unbacktraced leaves a "breadcrumb trail"
  if untried[s'] is empty then
      if unbacktracked[s'] is empty then return stop
      else a \leftarrow an action b such that result[s', b] = POP(unbacktracked[s'])
  else a \leftarrow POP(untried[s'])
  s \leftarrow s'
  return a
```



Important concepts that you should be able to explain and use now...

- Difference between the solution types:
 - a fixed actions sequence, and
 - a conditional plan (also called strategy or policy).
- What are belief states?
- How actions can be used to coerce the world into states.
- State estimation with repeated predict and update steps.
- The use of AND-OR trees.